




Separation of Concerns at Zero Cost

Modular Physics in C++20

Jędrzej Michalczyk

About me


- **Jędrzej Michalczyk** — R&D Engineer
- **SpaceForest** (Gdynia, Poland)
 - PERUN — single-stage suborbital launch vehicle
 - guidance, navigation & control software in modern C++
- The code in this talk is roughly :
github.com/jedrzejmichalczyk/accu-on-sea-2026-modular-physics



The original problem

- Simulate a rocket: **6-DOF dynamics + aerodynamics + propulsion + atmosphere + control**
- The same physics must run in three places:
 - engineering analysis (parameter studies, Monte Carlo)
 - hardware-in-the-loop test benches — real time
 - onboard — guidance runs against the same models, in real time
- **Correctness and performance are crucial**
- And: aero engineers, control engineers and software engineers all touch the same code

From symbolic math to embedded C



Sometimes it is an overkill.

01

The Principle

Separation of concerns in numerical code

≈ 5 min

THE MODELING PROBLEM

Where do the software modules go?



It is not obvious where to draw the software module boundaries.

*General good idea:
Choose correct states and make everything else a pure
state function.*

SOPOT — “Obviously Physics, Obviously Templates” — tries to capture this idea

What numerical code usually looks like

```
// derivs.cpp – all of physics in one function
void derivs(double t, const double* y, double* dydt) {
    // y[0..2] pos, y[3..5] vel, y[6..9] quat, y[10..12] omega, y[13] mass
    double rho = 1.225 * exp(-y[2] / 8500.0);           // atmosphere
    double v2   = y[3]*y[3] + y[4]*y[4] + y[5]*y[5];
    double Fd   = 0.5 * rho * v2 * CD * AREA;          // aero drag
    double Ft   = thrust_table_lookup(t);              // propulsion
    dydt[3] = (Ft*ax(y) - Fd*vx_hat(y)) / y[13];      // dynamics
    /* ... 400 more lines, every subsystem interleaved ... */
}
```


02

The Tradeoff That Isn't

Modularity vs performance — or so we're told

≈ 7 min

Attempt: classic OO interfaces

```
struct Component {  
    virtual void computeDerivatives(double t,  
                                    const double* state,  
                                    double* dydt) = 0;  
    virtual double getOutput(const std::string& name) = 0; // !!  
    virtual ~Component() = default;  
};  
  
std::vector<std::unique_ptr<Component>> components;  
for (auto& c : components) // innermost loop of RK4,  
    c->computeDerivatives(t, y, dydt); // 4x per step, 1000s of steps
```

PICK ONE?

The apparent tradeoff

Monolith

fast ✓

unmaintainable ✗

Virtual components

modular ✓

indirection in the hot loop ✗

Can be done


03

The Coupled Oscillator

Two masses, one spring, four components

≈ 8 min

THE RUNNING EXAMPLE



Who owns what — who asks what

Component	Owns (state)	Provides	Queries
Mass1 / Mass2	[position, velocity]	Position, Velocity, Mass	Force (whoever provides it)
Spring	— nothing	Force, Extension, PE	Position of both masses
EnergyMonitor	— nothing	TotalEnergy, Momentum	everything above

- Composition of the whole system
- Wiring
-

04

Code Walkthrough

The four components, for real

≈ 12 min

Tags: typed names for physical quantities

```
struct mass1 {  
    struct Position {};           // empty types – pure compile-time keys  
    struct Velocity {};  
    struct Force    {};  
    struct Mass     {};  
};  
  
struct mass2 { /* same shape, different types */ };  
  
namespace spring {  
    struct Extension    {};  
    struct PotentialEnergy {};  
};
```


PointMass: owns its state, asks for force

```
template<typename TagSet, Scalar T = double>
class PointMass final : public Component<2, T> {    // 2 states: [x, v]
    double m_mass;
public:
    template<typename Registry>
    LocalDerivative derivatives(T t, std::span<const T> state, const Registry& registry) const {
        T velocity = this->localState(state, 1);    // dx/dt = v
        T force = query<typename TagSet::Force>(registry, state);
        return {velocity, force / T(m_mass)};      // dv/dt = F/m
    }
    template<typename Registry>                    // Velocity: same, localState(s, 1)
    T compute(typename TagSet::Position, std::span<const T> s,
              const Registry&) const { return this->localState(s, 0); }
};
```

Spring: no state, just physics

```
template<typename TagSet1, typename TagSet2, Scalar T = double>
class Spring final : public Component<0, T> {           // 0 states!
    double m_stiffness, m_rest_length;
public:
    template<typename Registry>
    T compute(typename TagSet1::Force, std::span<const T> state, const Registry& reg) const {
        T x1 = query<typename TagSet1::Position>(reg, state);
        T x2 = query<typename TagSet2::Position>(reg, state);
        T extension = x1 - x2 - T(m_rest_length);
        return -T(m_stiffness) * extension;             // Hooke's law
    }
    template<typename Registry>
    T compute(typename TagSet2::Force, std::span<const T> state, const Registry& reg) const {
        return -query<typename TagSet1::Force>(reg, state); // Newton's 3rd
    }
};
```

EnergyMonitor: system totals from component data

```
template<Scalar T = double>
class EnergyMonitor final : public Component<0, T> {
public:
    template<typename Registry>
    T compute(system::TotalEnergy,
              std::span<const T> state, const Registry& registry) const {
        T m1 = query<mass1::Mass>(registry, state);
        T v1 = query<mass1::Velocity>(registry, state);
        T m2 = query<mass2::Mass>(registry, state);
        T v2 = query<mass2::Velocity>(registry, state);
        T pe = query<spring::PotentialEnergy>(registry, state);
        return T(0.5)*m1*v1*v1 + T(0.5)*m2*v2*v2 + pe;
    }
};
```

Wiring it together

```
auto sys = makeODESystem<double>(
    Mass1<double>(1.0, /*x0*/ 0.0, /*v0*/ 0.0),
    Mass2<double>(1.0, /*x0*/ 2.5, /*v0*/ 0.0),
    Spring12<double>(/*k*/ 4.0, /*L0*/ 1.0),
    EnergyMonitor<double>()
);


// the wiring is CHECKED at compile time:
static_assert(decltype(sys)::hasFunction<system::TotalEnergy>());

RK4Solver solver;
auto sol = solver.solve(derivs(sys), sys.getStateDimension(),
                        0.0, 20.0, /*dt*/ 0.001, sys.getInitialState());

double E = sys.query<system::TotalEnergy>(sol.states.back());
```

PROOF, NOT PROMISE

Does it work?



20 001

RK4 steps (dt = 1 ms)

16.7 ms

wall clock, g++ 14 -O2

1.4×10^{-13}

max relative energy drift over 20 s

05

How It Works

Tags, registry, concepts — no magic

≈ 10 min

Idea 1 — dispatch on type, not through a hierarchy

```
// instead of:
double getOutput(
    const std::string& name);

// one virtual + string compare
// per query, at runtime

// we write:
T compute(mass1::Position,
    std::span<const T> state) const;

// one overload per quantity,
// chosen by the COMPILER
```

- **Overload resolution replaces virtual dispatch**
- The tag argument is empty — it exists only to select the overload
- Unknown tag → compile error, with the tag's name in the message
- This is ordinary C++ — tag dispatch is as old as the STL (`std::advance`)

Idea 2 — the registry: 'who provides what'

```
template<typename T, ComponentConcept... Components>
class Registry {
    std::tuple<const Components&...> m_components;

    template<StateTagConcept Tag>
    static constexpr bool hasFunction() {
        return (ProvidesStateFunction<Components, Tag, T, Registry> || ...);
    }
                                     // fold over all components

    template<StateTagConcept Tag>
    auto computeFunction(std::span<const T> state) const {
        static_assert(hasFunction<Tag>(),
                       "No component provides this state function");
        const auto& provider = findProvider<Tag>(); // index found constexpr
        return provider.compute(Tag{}, state, *this);
    }
};
```


compute vs query: provide vs ask

```
// a component PROVIDES a quantity (one overload per tag):  
T compute(mass1::Force, std::span<const T> s, const Registry&) const;  
  
// anyone ASKS for one, by tag – component or driver:  
T f = query<mass1::Force>(registry, state);  
  
// query is just the registry call with friendlier syntax:  
template<StateTagConcept Tag, typename Registry, typename T>  
auto query(const Registry& reg, std::span<const T> s) {  
    return reg.computeFunction<Tag>(s);    // find provider -> provider.compute(...)  
}
```

Idea 3 — concepts make the contracts checkable

```
// what a component IS:
template<typename C>
concept ComponentConcept = requires(const C& c) {
    { C::state_size } -> std::convertible_to<size_t>;
    { c.getInitialLocalState() } -> std::same_as<typename C::LocalState>;
};

// what it PROVIDES (a quantity, selected by tag):
template<typename C, typename Tag, typename T, typename Registry>
concept ProvidesStateFunction =
    ComponentConcept<C> && StateTagConcept<Tag> &&
    requires(const C& c, std::span<const T> s, const Registry& reg)
    { c.compute(Tag{}, s, reg); };
```

Idea 4 — everything is a value

- The system is `std::tuple<Mass1, Mass2, Spring, EnergyMonitor>`
 - no heap, no pointers, no type erasure, no `std::function`
- **The optimizer sees the whole object graph as one struct**
 - every `compute()` call is a direct (and inlinable) call
- State lives in one contiguous vector; components know their offsets
- The price: component set is fixed at compile time
 - for flight software that's a feature — the vehicle doesn't grow a fin mid-flight

06

What the Compiler Sees

Trust, but disassemble

≈ 6 min

The abstraction compiles away

```

energy_sopot(span<const double>):      ; 4 components, registry dispatch
    mov    r10, [sys+272]
    movsd  xmm1, [.LC0]                ; 0.5
    mov    rax, [rcx]
    movsd  xmm0, [r10+8]               ; m1      <- runtime parameters,
    movsd  xmm3, [rax+rcx*8]           ; v1      loaded once
    mulsd  xmm0, xmm1
    mulsd  xmm0, xmm3                  ; 0.5*m1*v1
    mulsd  xmm0, xmm3                  ; ... * v1
    subsd  xmm2, [r8+16]               ; x1-x2-L0
    mulsd  xmm1, xmm2                  ; PE terms
    addsd  xmm0, xmm3
    addsd  xmm0, xmm1                  ; KE1+KE2+PE
    ret

```

0 calls · 0 jumps · 0 vtables — straight-line FP, the same shape as the hand-written formula.

...and the derivative? Same story.

```

mass1_derivative(...):
    mov    r9, [sys+112]
    mov    r10, [sys+128]
    movsd  xmm2, [r9+8]
    xorpd  xmm2, [.LC0]
    movsd  xmm0, [r8+rcx*8]
    subsd  xmm0, [r8+rdx*8]
    subsd  xmm0, [r9+16]
    movsd  xmm1, [r8+rcx*8]
    mulsd  xmm0, xmm2
    comisd  xmm3, 0
    jbe    .Lskip
    divsd  xmm0, [r10+8]
    unpcklpd xmm1, xmm0
    movups [rax], xmm1
    ret

```

; returns { dx/dt , dv/dt }
 ; spring component
 ; mass1 component
 ; stiffness k, negated -> -k
 ; x1
 ; - x2
 ; - L0 -> extension
 ; v1 (= dx/dt)
 ; -k * extension (spring force)
 ; if (damping > 0) - c == 0, so
 ; the branch is skipped
 ; force / m1 (= dv/dt)
 ; pack { v1, a1 }

the whole query chain — spring force, then $\div$ mass — is ~15 instructions, no calls.

MEASURED, NOT ESTIMATED

Numbers

~1 ns per state-function query — fully inlined

16.7 ms 20 001 RK4 steps of the 4-component oscillator (laptop, -O2)

1000× faster than an equivalent OO C# version

Why C++20 — and not C++26?

- **Everything here needs only C++20:**
 - concepts (the diagnostics!), fold expressions, `std::span`, CTAD
- **C++26 reflection (P2996) would remove the remaining boilerplate:**
 - tag declarations and `compute()` registration could be generated
 - we have working experiments in the repo: [reflect/](#)
- But flight software ships on the toolchain you have, **not the one you want**
 - embedded / cross targets lag years behind; certification lags more


07

Where It Scales

From two masses to a flying vehicle

≈ 7 min

Same pattern, real vehicle: 6-DOF rocket



- **11 components, 14 states** — position, velocity, quaternion, angular rate, mass
- Thrust vector control & attitude determination: same tags-and-registry architecture
- **These patterns run in production flight software** — correctness and performance both non-negotiable

Bonus: the scalar is a template parameter

```
// simulate:
auto sim = makeODESystem<double>(comps...);

// differentiate — SAME components, different scalar:
using D = Dual<double, 14>; // forward-mode autodiff
auto grad = makeODESystem<D>(comps...);
auto J = computeJacobian(grad, t, state); // exact, to machine precision

// linearize -> LQR gain straight from the nonlinear model
auto K = lqr(J.A, J.B, Q, R);
```

- **Jacobians without hand-derived math or finite differences**
- The controller you'll see in a minute was designed exactly this way

Can we go further?

- **Frames as types** — body vs ENU vs ECEF: mixing them up should be a compile error too
- **Physical laws as concepts** — 'conserves energy', 'action-reaction pair' as checkable contracts
- Units in the type system — dimensional analysis at compile time

The goal, all the way down: a consistent level of abstraction

We've been here before

[PLACEHOLDER — scan of rocket-control code from a classic textbook
(Zipfel / 1960s–90s Fortran, zero modularity)]

- Six decades of rocket code — the math was always modular; the code never was
- **Not because the engineers were careless — because the languages made them choose**

One more thing: six pendulums on a cart

- **Inverted 6-pendulum chain** balanced on a moving cart — LQR-stabilized
- Built from the same components, tags and registry you just saw
- Linearized with `Dual<double, N>` — the autodiff Jacobian from two slides ago
- Same C++20, compiled to WebAssembly, running in your browser right now:

jedrzejmichalczyk.github.io/sopot

▶ LIVE DEMO

Takeaways

- **Separation of concerns is achievable at zero cost** in C++20 — measured, not asserted
- The recipe is small: **typed tags + overloaded compute() + a constexpr registry + concepts**
 - no inheritance, no virtuals, no strings, no codegen
- **Prefer abstractions you can verify** — conservation laws and disassembly are unit tests for architecture
- **Quality vs performance was the false tradeoff** — modern C++ finally lets us refuse it

Thank you! Questions?

- Code & slides: github.com/jedrzejmichalczyk/sopot
- Live demo: jedrzejmichalczyk.github.io/sopot
- Email: jedrzej.michalczyk@gmail.com